

ECHOES: EQUAL-WEIGHT COMPLETED HYPOTHETICAL OBSERVATION ENSEMBLES FOR BOSS DR12 CMASS-SOUTH: POSTERIOR COMPLETED-CATALOG SAMPLES FOR GALAXY CLUSTERING BEYOND TWO POINTS

TOM ABEL

Kavli Institute for Particle Astrophysics and Cosmology, Stanford University, Stanford, CA, USA

Draft version June 18, 2026

ABSTRACT

Most precision galaxy-clustering analyses communicate between survey teams and theory teams through a small set of summary statistics, most often two-point correlation functions or power spectra. This convention is effective because the observational effects that enter these statistics have been studied for decades, but it also makes each new statistic responsible for reinterpreting the survey window, fiber assignment, redshift failures, imaging systematics, and random catalogs. We present ECHOES, Equal-weight Completed Hypothetical Observation Ensembles, for BOSS DR12 CMASS-South: an ensemble of equal-weight, cosmology-free completed catalogs that keep every observed galaxy at its measured (α, δ, z) and restore the spectroscopically missing galaxies at their real SDSS imaging positions. The unknown redshifts are sampled from a data-driven posterior combining a color-space k -nearest-neighbor photo- z , an empirical close-pair prior, and a local-density line-of-sight prior. Smooth imaging-systematic weights are represented by local analog galaxies rather than by per-object completeness weights. The released posterior is compact: one 2.04 MB file stores the fixed observed catalog and the inverse CDF for each restored target, and a numpy-only sampler generates reproducible complete catalogs from integer seeds. For seed 0 the sample contains 119,923 equal-weight galaxies, of which 109,636 are observed, 5,272 are restored fiber-collision targets, 1,505 are restored redshift failures, and 3,510 are imaging-systematic analogs. We validate the product against truth-known completions and MultiDark-Patchy mocks, finding that projected clustering is recovered to the 1–2% level over $0.5 < r_p < 40 h^{-1} \text{Mpc}$, while the posterior completion uncertainty in $w_p(r_p)$ is only 0.17–0.44% per bin, below the sample-variance scale of the BOSS South Galactic Cap. We also describe a fully complementary graphGP route: a conditional anisotropic Gaussian-process posterior over the galaxy density field, sampled by Matheron’s rule on a sparse graph and evaluated along each missing galaxy’s line of sight. The resulting catalogs are intended to let analysts run two-point, higher-order, marked, nearest-neighbor, topological, field-level, or simulation-based statistics directly on completed catalogs while propagating the observational-completion uncertainty by resampling the catalog posterior.

Subject headings: Large-scale structure of the universe (902); Galaxy surveys (598); Astrostatistics (1882); Cosmological parameters (339); Observational cosmology (1146)

1. INTRODUCTION

The central compression in observational large-scale structure is not a simulation or an estimator, but an agreement about what object is being communicated. A survey team measures a summary statistic from a complicated instrumental and astrophysical selection function; a theory team predicts the same statistic for a cosmological and galaxy-formation model. The summary statistic is the interface. For modern spectroscopic galaxy surveys that interface has usually been a two-point statistic: a correlation function, a projected correlation function, a redshift-space multipole, or a Fourier-space power spectrum (M. Davis & P. J. E. Peebles 1983; S. D. Landy & A. S. Szalay 1993; A. J. S. Hamilton 1993; H. A. Feldman et al. 1994). This choice is not merely historical. The observational corrections for two-point clustering have been revisited repeatedly, calibrated in mocks, and encoded in survey catalogs, random catalogs, and weights.

At the same time, the cosmological information in a galaxy survey is not exhausted by two-point clustering.

The literature contains a long sequence of proposed and demonstrated summaries: higher-order moments and hierarchical statistics (I. Szapudi & A. S. Szalay 1993; J. N. Fry & E. Gaztanaga 1993), void-probability functions (J. Einasto et al. 1991), counts-in-cells and one-point density PDFs (C. M. Baugh et al. 1995; I. Szapudi et al. 1996; C. Uhlemann et al. 2020; O. Friedrich et al. 2020), marked correlation functions and marked power spectra (G. Harker et al. 2006; S. Satpathy et al. 2019; E. Massara et al. 2021, 2023), nearest-neighbor distributions (A. Banerjee & T. Abel 2021; S. Yuan et al. 2024), wavelet and scattering-transform statistics (S. Cheng et al. 2020; G. Valogiannis & C. Dvorkin 2022; B. Régalo-Saint Blancard et al. 2024), persistent-homology summaries of the cosmic web (G. Wilding et al. 2021; M. Biagetti et al. 2022; J. H. T. Yip et al. 2024), and field-level or simulation-based inference methods (J. Jasche & B. D. Wandelt 2013; N.-M. Nguyen et al. 2024; P. Lemos et al. 2024; C. Hahn et al. 2023a,b, 2024a,b; E. Massara et al. 2025). Large simulation suites such as Quijote (F. Villaescusa-Navarro et al. 2020) have been essential for calibrating and stress-testing many of these

summaries. The Beyond-2pt Collaboration’s parameter-masked mock-data challenge (ThBeyond-2pt Collaboration et al. 2025) makes the same point in a deliberately broad form: once the Gaussian, two-point description is relaxed, many summaries can extract complementary information, and their relative utility is an empirical question.

The barrier between those ideas and real spectroscopic data is often not the statistic itself. It is the observational model. A statistic that is easy to define on a periodic simulation box can become difficult to interpret in a survey with angular holes, variable targeting, fiber collisions, redshift failures, stellar density, seeing, extinction, and redshift-dependent selection. For two-point statistics, decades of work have made the corresponding correction pathways usable. For many newer statistics, the same corrections are either rederived from scratch, approximated, or avoided by restricting the analysis to mocks and forecasts. This paper asks whether the interface can be moved one step earlier: instead of releasing only weighted galaxies and randoms and asking every statistic to understand the weights, release a posterior over completed galaxy catalogs that have already absorbed the main observational incompleteness into equal-weight point sets.

This perspective has a long historical lineage. The first large maps of extragalactic structure already had to confront the fact that the sky is not a uniform observing surface. E. Hubble (1934) discussed the distribution of extragalactic nebulae in the presence of Galactic obscuration, a problem that later became a quantitative dust-correction program culminating in all-sky reddening maps such as D. J. Schlegel et al. (1998). Redshift surveys then turned angular inhomogeneity into a three-dimensional selection problem (M. Davis & P. J. E. Peebles 1983; V. de Lapparent et al. 1986). The 2dF Galaxy Redshift Survey and SDSS established the modern industrial mode of galaxy redshift surveys (D. G. York et al. 2000; M. Colless et al. 2001; S. Cole et al. 2005), and the SDSS luminous red galaxy sample made the two-point correlation function a precision cosmological observable with the detection of the baryon acoustic feature (D. J. Eisenstein et al. 2005).

BOSS inherited this history and made it operational. Its spectrographs, fiber system, target selection, and large-scale-structure catalogs are documented in K. S. Dawson et al. (2013); S. A. Smee et al. (2013); B. Reid et al. (2016). The final DR12 cosmology analyses used a mature system of angular masks, radial random catalogs, redshift-failure weights, close-pair weights, imaging-systematic weights, and FKP weights (S. Alam et al. 2017; A. J. Ross et al. 2017). The resulting measurements were robust because the effect of these ingredients on the two-point statistics had been studied. But the same survey artifacts have no reason to couple only to two-point functions. Fiber collisions remove close angular pairs; redshift failures correlate with observing conditions and target properties; imaging systematics modulate the target density across degree scales; and the random catalog carries a particular representation of the angular and radial selection function.

The same issues will not become simpler in later surveys. DESI’s robotic fiber-positioner system and its fiber-assignment algorithm create a richer assignment

problem than the BOSS plug-plate geometry (J. H. Silber et al. 2023; L. Pinol et al. 2017; T.-W. Lan et al. 2023; J. Lasker et al. 2025). Euclid, the Nancy Grace Roman Space Telescope, the Chinese Space Station Telescope, and the Vera C. Rubin Observatory’s LSST will produce enormous galaxy samples with their own mixtures of imaging, spectroscopy, selection, and calibration systematics (Euclid Collaboration et al. 2025; Y. Wang et al. 2022; H. Miao et al. 2023; CSST Collaboration et al. 2026; H. Zhan & J. A. Tyson 2018). If those surveys wish to serve a community that will analyze the data with statistics not yet standardized, the most useful data product may not be a single corrected summary statistic. It may be a completed probabilistic catalog.

This work builds such a product for BOSS DR12 CMASS-South. We call the approach ECHOES: Equal-weight Completed Hypothetical Observation Ensembles, echoes of the survey one would have taken if the principal spectroscopic incompleteness and smooth imaging systematics were negligible. It is not a new cosmological measurement. It is a data-modeling and validation paper. We construct an ensemble of completed catalogs conditioned on the published BOSS data and SDSS imaging: observed spectroscopic galaxies remain fixed, missing targets are restored at real imaging positions, and redshift uncertainty is sampled. The ensemble is cosmology-free in the sense that it contains only RA, Dec, and redshift; an analyst chooses a fiducial cosmology only when measuring three-dimensional clustering. The ECHOES product is designed so that the summary statistic, rather than the observational correction, can once again be the object under study.

2. ECHOES COMPLETED-SURVEY METHODOLOGY

2.1. Data model and scope

The target ECHOES product is an ensemble

$$\mathcal{C}^{(s)} = \{\alpha_i^{(s)}, \delta_i^{(s)}, z_i^{(s)}, p_i^{(s)}\}_{i=1}^{N_s}, \quad (1)$$

where s is a seed and p_i is a provenance code. The coordinates are observed RA, Dec, and redshift. No comoving coordinate, Alcock–Paczynski scaling (C. Alcock & B. Paczynski 1979), redshift-space model (N. Kaiser 1987), or fiducial cosmology is stored in the catalog. The completion posterior is conditional on the observed BOSS catalog, the SDSS imaging target catalog, and the BOSS angular mask and weights. It does not represent cosmic variance; it represents the uncertainty in replacing spectroscopic incompleteness and smooth imaging-systematic weights by equal-weight galaxies.

The ECHOES product corrects three effects in the default release: fiber collisions, redshift failures, and the smooth imaging-systematic modulation represented by `WEIGHT_SYSTOT`. An optional inpainted version also fills interior mask holes, but the default scientific product leaves the survey mask explicit. The FKP weight (H. A. Feldman et al. 1994) is not a completeness correction and is not converted into galaxies. Analysts may still apply FKP or other inverse-variance weights when estimating a statistic; the released point catalog itself is equal weight.

2.2. BOSS CMASS-South sample

We use the BOSS DR12 CMASS South Galactic Cap subset with the same clean cuts used in several modern analyses of these data:

$$\text{RA} < 28^\circ \text{ or } \text{RA} > 335^\circ, \quad \text{Dec} > -6^\circ, \quad 0.45 < z < 0.60. \quad (2)$$

The input contains 109,636 observed spectroscopic galaxies. The BOSS completeness weights imply a mean completeness correction of 1.074, or about 7.4% missing objects before the smooth imaging-systematic analog step. In the real target recovery used here, 5,272 fiber-collision targets and 1,505 redshift-failure targets are restored, for 6,777 spectroscopically missing galaxies in the compact posterior package.

The BOSS loader reads the standard large-scale-structure columns `RA`, `DEC`, `Z`, `WEIGHT_SYSTOT`, `WEIGHT_NOZ`, `WEIGHT_CP`, and `WEIGHT_FKP`. It also reads `MODELFLUX`, `EXTINCTION`, `IMATCH`, and `ICOLLIDED` when photometry is requested. Extinction-corrected Pogson magnitudes are computed as

$$m_b = 22.5 - 2.5 \log_{10} f_b - A_b, \quad (3)$$

with adjacent colors $(u - g, g - r, r - i, i - z)$. The photometric-redshift model uses $(g - r, r - i, i - z, i)$, dropping the u band because it is noisy for the red CMASS population.

2.3. Recovering the missing target list

The completion uses real imaging positions for missing galaxies. We start from the DR12 CMASS color-selected target catalog, cross-matched to the good redshift BOSS LSS catalog. A target with a spectrum but nonzero `ZWARNING` is classified as a redshift failure. A target with no usable spectrum is eligible to be a fiber-collision target. Because the no-spectrum photometric pool also contains objects that were not part of the final BOSS LSS sample, it is tied back to the BOSS book-keeping. For each observed survivor with `WEIGHT_CP` > 1, the algorithm greedily claims $\text{round}(\text{WEIGHT_CP} - 1)$ nearest unclaimed no-spectrum targets within the BOSS fiber-collision scale, 62". Redshift failures are tied to survivors with `WEIGHT_NOZ` > 1 using the analogous $\text{round}(\text{WEIGHT_NOZ} - 1)$ count and a wider 15' neighborhood. This preserves the integer count implied by the official weights while placing the restored galaxies at their actual imaging positions and colors.

2.4. Color-space photo- z posterior

For each missing target we require a redshift posterior, not only a point estimate. We fit a nonparametric k -nearest-neighbor model to the observed CMASS galaxies in the whitened four-dimensional feature space $(g - r, r - i, i - z, i)$. The default model uses $k = 100$ and a `ckdTree`. For a query feature vector $\mathbf{x}$, the posterior support is the set of neighbor redshifts $\{z_j\}_{j=1}^k$ with inverse-distance weights

$$w_j(\mathbf{x}) = \frac{(d_j^2 + \epsilon)^{-1}}{\sum_{\ell=1}^k (d_\ell^2 + \epsilon)^{-1}}. \quad (4)$$

The held-out diagnostic in the report gives $\sigma_{\text{NMAD}} = 0.019$, a bias of -0.0006 , a catastrophic outlier rate of 1.8%, and a probability-integral-transform mean of 0.503. These numbers are adequate for the role the

photo- z plays here: it supplies a color likelihood that is then multiplied by a local-density prior along the same line of sight.

2.5. Local-density redshift posterior

The default redshift engine is a fast local-density approximation to a conditional field posterior. For missing target i at angular position $\hat{\mathbf{n}}_i$ and colors $\mathbf{x}_i$, redshifts are sampled from

$$p_i(z) \propto \hat{\lambda}_{\text{loc}}(\hat{\mathbf{n}}_i, z) \hat{p}_{\text{col}}(z | \mathbf{x}_i) q_i(z), \quad (5)$$

where $\hat{p}_{\text{col}}$ is the color-space photo- z density, q_i is an empirical close-pair prior for collisions and unity for redshift failures, and $\hat{\lambda}_{\text{loc}}$ is the observed local line-of-sight density. The local density is estimated from the $K = 150$ nearest observed spectroscopic galaxies in angular unit-vector space:

$$\hat{\lambda}_{\text{loc}}(\hat{\mathbf{n}}_i, z) = \sum_{j \in \mathcal{N}_K(i)} \exp \left[-\frac{(z - z_j)^2}{2\sigma_f^2} \right], \quad \sigma_f = 0.004. \quad (6)$$

The photo- z factor is evaluated on the same 256-point redshift grid with a Gaussian smoothing scale $\sigma_p = 0.02$ applied to the neighbor redshifts and their k NN weights. For collided targets, $q_i(z)$ is the symmetrized histogram of redshift differences among observed close pairs within 62", measured directly from the data over $|\Delta z| < 0.06$. The sampled redshift is then jittered by a Gaussian of width $\sigma_f/2$ and clipped to the CMASS redshift interval.

Equation (5) is deliberately observational. It does not assign a missing galaxy to a simulated halo or require a fiducial cosmology. The object is placed at its real angular position and on a redshift structure seen in the neighboring spectroscopy, with the color likelihood preventing the local-density prior from assigning obviously incompatible redshifts.

2.6. Imaging-systematic analogs and provenance

The BOSS imaging-systematic correction `WEIGHT_SYSTOT` is a smooth density correction, not a statement that a particular nearby target failed to receive a fiber. We therefore do not convert it into exact duplicate galaxies. For every observed or spectroscopically restored base object, the sampler adds

$$n_{\text{extra}} = \lfloor \max(w_{\text{systot}} - 1, 0) + U \rfloor, \quad U \sim \mathcal{U}(0, 1), \quad (7)$$

local analogs at the same source position with a 1" Gaussian angular jitter and the same redshift. This preserves resolved angular and three-dimensional clustering while avoiding a zero-separation spike that would corrupt nearest-neighbor or counts-in-cells statistics. The default implementation keeps all real detections; it restores density deficits implied by $w_{\text{systot}} > 1$ rather than thinning observed galaxies in regions with $w_{\text{systot}} < 1$.

Each output object carries a provenance code: 0 for observed spectroscopy, 1 for restored fiber-collision target, 2 for restored redshift failure, 3 for imaging-systematic analog, 4 for rare host-redshift fallback, and 5 for optional mask inpainting. The provenance vector is central to reproducibility because analysts can rerun a statistic on individual components or test sensitivity to the imaging analog prescription.




FIG. 1.— Inputs to the completion. *Left*: the BOSS DR12 CMASS-South footprint and observed-galaxy distribution on the sky. *Right*: the published completeness weights — fiber-collision w_{CP} , redshift-failure w_{NOZ} , and imaging-systematic w_{SYSTOT} — whose product defines the per-region incompleteness the completion restores. The completion is conditioned on this geometry, on these weights, and on the SDSS imaging positions of the spectroscopically missing targets; it converts the weights into explicit equal-weight galaxies rather than carrying them as per-object multipliers.




FIG. 2.— The color-space k -nearest-neighbor photo- z that supplies the color likelihood in the redshift posterior, validated on a held-out sample of good-redshift CMASS galaxies. *Left*: posterior-median photo- z versus spectroscopic redshift ($\sigma_{\text{NMAD}} \simeq 0.019$). *Middle*: the probability-integral-transform histogram — the rank of each true redshift within its own posterior — which is flat (mean $\simeq 0.5$), so the posterior is calibrated and can be sampled faithfully. *Right*: a single posterior draw per object reproduces the held-out $n(z)$. The photo- z enters the completion only as the likelihood factor $\hat{p}_{\text{col}}(z | \mathbf{x})$, never as a standalone redshift.

2.7. Compact posterior and sampling

The default posterior has a simple product structure once the observed data are fixed. The observed galaxies never change across realizations. The missing galaxy positions never change. The redshift posterior for each missing galaxy, Equation (5), is also fixed after it is computed. A realization is only a draw from each per-object inverse CDF plus the stochastic imaging analogs.

The release therefore stores this posterior as a single compressed file of 2.04 MB: the fixed observed catalog, the base positions and provenance for the observed and restored targets, `WEIGHT_SYSTOT`, and a tabulated inverse-CDF redshift posterior for each of the 6,777 missing targets (the file schema and quantization are given in Appendix C). A matching uniform-footprint random catalog contains 438,544 points; because CMASS-South is approximately 99% complete in the angular mask, these randoms are uniform over the footprint to the percent level.

Sampling is deterministic under a seed:

1. draw one uniform variate for each missing target and invert its stored redshift CDF;
2. concatenate the fixed observed galaxies and the restored targets;
3. draw local analog counts from `WEIGHT_SYSTOT`;
4. return arrays `ra`, `dec`, `z`, `prov`, and `N`.

For seed 0, the sampled catalog contains 119,923 objects: 109,636 observed, 5,272 collision restorations, 1,505 redshift-failure restorations, and 3,510 smooth-systematic analogs. Without the imaging analogs, the fixed base catalog contains 116,413 objects.

2.8. Clustering measurements and acceptance metrics

For angular clustering we measure Landy–Szalay estimates (S. D. Landy & A. S. Szalay 1993) with the fixed random catalog. For three-dimensional statistics we use `Corrfunc` (M. Sinha & L. H. Garrison 2020). The catalog itself is cosmology-free; `Corrfunc` measurements convert redshift to comoving distance at measurement time, using the analyst’s fiducial cosmology, and then call the OpenMP mock pair counters with `is_comoving_dist=True`. The main validation statistics are $w(\theta)$, $w_p(r_p)$, $\xi(r_p, \pi)$, and redshift-space multipoles $\xi_0(s)$ and $\xi_2(s)$. We also evaluate higher-order or non-two-point summaries: nearest-neighbor CDFs and counts-in-cells. These are deliberately included because the purpose of the catalog posterior is to support statistics beyond those for which BOSS weights were optimized.

The acceptance criteria are empirical. In truth-known tests, the completed ensemble should recover the truth catalog to the percent level over the scales where the input mocks themselves are faithful. On the real data, the equal-weight completed catalogs should reproduce the official $w_c = \text{WEIGHT_SYSTOT}(\text{WEIGHT_CP} + \text{WEIGHT_NOZ} - 1)$ weighted BOSS clustering within the known scale of the correction. The completion posterior spread should be smaller than the sample-variance covariance for BOSS CMASS-South, because it represents only uncertainty in observational completion conditioned on the observed catalog.

3. THE GRAPHGP COMPLEMENT

The local-density engine in Section 2.5 is fast and compressible, but it treats the missing-galaxy redshift posteriors as conditionally independent once the observed catalog is fixed. The graphGP route is a complementary, more explicit field-level model. It assumes that the complete galaxy overdensity field can be approximated by a Gaussian process with covariance tied to the observed two-point function, and that the observed galaxies are noisy samples of that field.

Let $\delta(\mathbf{x})$ be the galaxy overdensity in a fiducial comoving coordinate system. The graphGP prior is

$$\delta(\mathbf{x}) \sim \mathcal{GP}[0, K_\xi(|\mathbf{x} - \mathbf{x}'|)], \quad (8)$$

where $K_\xi(r)$ is a tabulated covariance derived from a Landy–Szalay measurement of the galaxy correlation function. The likelihood is a Gaussian approximation to an inhomogeneous Poisson process. At observed galaxy




FIG. 3.— Restored missing targets and the resulting redshift distribution. *Left*: a zoomed sky region showing observed galaxies together with the recovered fiber-collision and redshift-failure targets, each placed at its real SDSS imaging position; only the redshift of each restored galaxy is uncertain. *Right*: the equal-weight completed $n(z)$ reproduces the completeness-weighted observed $n(z)$, and the thin-slice panel illustrates that *which* restored galaxies fall in a given redshift slice varies between realizations — this realization-to-realization variation is the completion posterior. Realizations differ only in the sampled redshifts of the missing galaxies and in the stochastic imaging-systematic analogs.

locations $\mathbf{x}_D$, the data vector is

$$y_i = \bar{n}_i^{-1} - 1, \quad N_{ii} = \bar{n}_i^{-1}, \quad (9)$$

where $\bar{n}_i$ is the local mean density implied by the survey selection. The posterior is Gaussian:

$$\mu_{\text{post}}(\mathbf{x}) = K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} y_D, \quad (10)$$

$$C_{\text{post}}(\mathbf{x}, \mathbf{x}') = K(\mathbf{x}, \mathbf{x}') - K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} K(\mathbf{x}_D, \mathbf{x}'). \quad (11)$$

The implementation samples this posterior by Mathéron’s rule (G. B. Rybicki & W. H. Press 1992). A prior draw δ_{prior} is generated on a sparse neighbor graph using a Vecchia-style approximation (M. Katzfuss et al. 2020; F. Schäfer et al. 2021). The conditional draw is

$$\delta_{\text{post}}(\mathbf{x}) = \delta_{\text{prior}}(\mathbf{x}) + K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} [y_D - \delta_{\text{prior}}(\mathbf{x}_D)]. \quad (12)$$

The expensive solve is amortized over many missing-galaxy redshift draws. The primary stored object is a light-cone-native array $1 + \delta_{\text{post}}$ on redshift shells and HealPIX pixels. For a missing galaxy, graphGP replaces the local KDE in Equation (5) by the corresponding field draw:

$$p_i(z | s) \propto [1 + \delta_{\text{post}}^{(s)}(\hat{\mathbf{n}}_i, z)] \bar{n}(z) \hat{p}_{\text{col}}(z | \mathbf{x}_i) q_i(z). \quad (13)$$

All missing galaxies in realization s share the same density-field draw, so their redshift assignments are correlated.

This model has a fiducial comoving coordinate system internally because a three-dimensional covariance needs a distance. That coordinate system is a modeling gauge for the field interpolation, not a cosmology baked into the released KNN product. The report tests this explicitly by rebuilding the graphGP route with different fiducial distance-redshift relations. Planck-like and Einstein-de Sitter choices give per-object redshift RMS differences of about 0.001–0.0012, correlations of 0.9995–0.9996, and changes in $w_p(r_p)$ below 0.1%. Within the accuracy demanded by the present BOSS completion, graphGP is therefore best viewed as a correlated posterior engine whose output is stable to reasonable internal coordinate choices.

4. APPLICATION TO BOSS DR12 CMASS-SOUTH

4.1. Why BOSS is the right test case

BOSS DR12 is the ideal laboratory because the official two-point analyses are mature. The survey, target selection, and LSS catalogs are documented (K. S. Dawson et al. 2013; B. Reid et al. 2016), the spectroscopic system is understood (S. A. Smee et al. 2013), the DR12 cosmology and systematic-correction papers provide community reference measurements (S. Alam et al. 2017; A. J. Ross et al. 2017), and independent mock suites such as MultiDark-Patchy exist for validation (F.-S. Kitaura et al. 2016). A completed-catalog product can therefore be judged against a well-established reference: it should reproduce the weighted BOSS two-point measurements while exposing a point catalog suitable for statistics that were not part of the official BOSS pipeline.

The official BOSS correction strategy keeps the observed spectroscopic catalog and supplies per-object weights. Close-pair and redshift-failure weights move the missing object’s statistical weight to an observed neighbor, while `WEIGHT_SYSTOT` corrects smooth density modulations associated with imaging and observing conditions. Random catalogs encode the angular mask and radial selection. This system is extremely effective for the published two-point program, but for a general statistic the analyst must decide how to apply weights, whether the statistic is linear or nonlinear in those weights, and how random catalogs should enter. Fiber-collision corrections illustrate the point: nearest-neighbor upweighting and more refined pairwise or angular correction schemes can be calibrated for two-point functions (H. Guo et al. 2012; F. G. Mohammad et al. 2020), but they do not automatically define how a nonlinear statistic should respond to a missing close pair.

SIMBIG is a useful modern comparison point. It uses BOSS CMASS-South-like cuts and a forward-modeling program to connect nonlinear galaxy clustering statistics to cosmological parameters (C. Hahn et al. 2023a,b, 2024a,b). Later work extends the same style of analysis to summaries such as the wavelet scattering transform and marked power spectra (B. Régalo-Saint-Blancard et al. 2024; E. Massara et al. 2025). The scientific ambition is close to ours—to use information beyond the Gaussian two-point limit—but the observational interface remains the published BOSS selection and weighting sys-

tem. Our product is complementary: it provides an ensemble of completed data catalogs so the observational correction can be sampled once and then reused across many statistics.

4.2. Default KNN-field completion

The default BOSS ECHOES product uses the KNN-field engine. The report’s top-level numbers are as follows. The observed catalog has 109,636 CMASS-South galaxies. The real missing-target list contains 5,272 collision targets and 1,505 redshift-failure targets. The fixed base catalog therefore has 116,413 galaxies. With the imaging-systematic analog sampler enabled, seed 0 has 119,923 galaxies, including 3,510 analogs. Seeds 0–4 vary only at the $\simeq 0.1\%$ level in total count, as expected for the smooth analog step.

Figure 3 shows the restored-target population and sampled $n(z)$. Figure 4 compares the angular and two-dimensional clustering, and Figure 5 the projected correlation and redshift-space multipoles. The completed equal-weight catalogs reproduce the official weighted angular clustering to roughly 1.5% over the plotted angular range and the projected and multipole statistics at the 1–5% level, depending on scale and statistic. This is the expected accuracy target: the completed catalog is not an exact algebraic rewrite of the official weights, because it assigns redshifts and positions to missing galaxies rather than moving all their weight onto observed neighbors, but it should remain inside the systematic scale already tolerated by the reference BOSS treatment.

4.3. Truth-known validation

Two truth-known tests are used. The first hides the spectroscopic redshifts of a subset of real BOSS galaxies and then asks the completion machinery to recover the clustering of the original full catalog. The second applies the same machinery to independent MultiDark-Patchy mocks. In both tests, an “oracle” catalog that restores the hidden objects with their true redshifts provides a floor: if the oracle differs from truth, the difference comes from the masking and bookkeeping rather than from redshift assignment.

The main result is shown in Figure 6. The real-BOSS-truth mock and the Patchy mocks recover $w_p(r_p)$ to about 1–2% over $0.5 < r_p < 40 h^{-1}\text{Mpc}$. The oracle ratio lies in the range 0.997–1.005, so the remaining error is dominated by the redshift assignment posterior. The Patchy test is faithful above about $1 h^{-1}\text{Mpc}$; the sub-Mpc deficit is a known one-halo limitation of the Patchy mocks rather than a failure mode inferred from real BOSS data.

4.4. Higher-order and non-two-point checks

The validation intentionally includes summaries that are nonlinear in the point set. For the nearest-neighbor distance CDFs (A. Banerjee & T. Abel 2021), the maximum absolute CDF differences between completed and truth catalogs are 0.0022, 0.0021, and 0.0011 for $k = 1, 2, 4$, respectively. For the counts-in-cells distribution (C. Uhlemann et al. 2020; O. Friedrich et al. 2020) in $8 h^{-1}\text{Mpc}$ cells, the mean is 0.660 versus 0.663 in truth, the variance-to-mean ratio is 2.37 versus 2.42, and the skewness is 2.90 versus 2.94. These tests are not exhaustive, but they are important because a product that only

reproduces w_p could still fail for nearest-neighbor or cell-count statistics.

4.5. Calibration of completion uncertainty

The completion posterior is intentionally narrower than a full survey posterior. It conditions on the observed CMASS-South realization and resamples only the unobserved spectroscopic information and smooth imaging analogs. In the Patchy calibration suite, the posterior uncertainty in $w_p(r_p)$ is 0.17–0.44% per bin, while sample variance over the same bins is 0.8–7.2%. The completion therefore adds only 6–36% of the sample-variance error in quadrature, depending on scale. The 68% completion band brackets the mock truth in only about 8% of bins because it is not intended to regenerate the large-scale density field. This is correct behavior: a user should add the completion covariance to a cosmic-variance or simulation covariance, not replace that covariance with the completion ensemble.

4.6. KNN-field versus graphGP

The KNN-field and graphGP routes agree on the main scientific conclusion but make different compromises. The KNN-field engine is sharper on the sub-Mpc collision scale because it directly uses nearby observed redshifts and the empirical close-pair prior. graphGP is smoother and more explicitly correlated across missing objects because all lines of sight share a density-field draw. In truth recovery, both routes recover the input clustering to the percent level. In the real-data head-to-head comparison, graphGP lies closer to the official weighted reference on larger projected scales, with a typical ratio near unity, while the KNN-field realization is modestly high in some bins. Because the KNN-field route is much cheaper and compresses to the compact public ECHOES posterior, it is the default release. graphGP is retained as a first-class alternative for surveys or statistics where correlated redshift-completion uncertainty is essential.

Although the graphGP field is built in a fiducial comoving frame, that choice is a gauge rather than a prior. Repeating the redshift assignment under two extreme fiducial cosmologies — Planck ($\Omega_m = 0.315$) and Einstein-de Sitter ($\Omega_m = 1.0$) — changes the per-object redshifts by an RMS of only ~ 0.001 (correlation 0.9996), an order of magnitude below the intrinsic assignment scatter, and the recovered $w_p(r_p)$ by less than 0.1% (Figure 10). Because the kernel is measured from the data and the output is expressed in observed redshift, the graphGP route is data-driven and carries no cosmological prior.

5. SAMPLING THE COMPLETION POSTERIOR

For the default ECHOES release, a posterior sample is a completed catalog seed. The mean of a statistic over seeds estimates the statistic for the completed survey, and the covariance over seeds estimates uncertainty due to the completion. If $S[\mathcal{C}^{(s)}]$ is any statistic of a catalog,




FIG. 4.— Two-point clustering of the equal-weight completed catalogs on the real BOSS CMASS-South data, with no completeness weights applied. *Left*: the angular correlation function $w(\theta)$ of the completed ensemble (mean and realization scatter, blue band) against the official w_c -weighted measurement (orange); the two agree to $\simeq 1.5\%$ across 0.05° – 2° , and the ratio panel shows no scale-dependent trend. *Right*: the two-dimensional correlation $\xi(\Delta\theta, \Delta z)$ measured directly in observed coordinates and its completed/weighted ratio across the plane. The anisotropy between angular and line-of-sight separation is reproduced to a few percent, so the geometric (Alcock–Paczynski) information carried by the full $\xi(\Delta\theta, \Delta z)$ plane is preserved without imposing a distance–redshift relation. The completion replaces survey weights by explicit equal-weight galaxies while leaving the measured two-point clustering unchanged at the percent level.




FIG. 5.— Standard redshift-space clustering of the equal-weight completed ensemble (no completeness weights) versus the official w_c -weighted BOSS measurement, with a Planck fiducial cosmology used only to convert redshift to distance. The projected correlation $w_p(r_p)$ (with the completed/official ratio) and the redshift-space monopole and quadrupole ξ_0, ξ_2 agree with the weighted reference at the 1–5% level over $0.5 < r_p < 40 h^{-1} \text{Mpc}$. An equal-weight completed catalog is therefore a drop-in replacement for the weighted catalog in the standard two-point estimators.

then

$$\bar{S} = \frac{1}{N_s} \sum_{s=1}^{N_s} S[\mathcal{C}^{(s)}], \quad (14)$$

$$C_{ab}^{\text{comp}} = \frac{1}{N_s - 1} \sum_s \left(S_a[\mathcal{C}^{(s)}] - \bar{S}_a \right) \left(S_b[\mathcal{C}^{(s)}] - \bar{S}_b \right). \quad (15)$$

This covariance should be added to, or modeled jointly with, the ordinary survey covariance. It should not be interpreted as the covariance of the cosmological density field.

For graphGP the same equations apply, but the meaning of a seed changes. A seed selects a conditional density-field draw as well as the missing-galaxy redshift draws from Equation (13). The redshift assignments are therefore correlated across missing galaxies. This is the more faithful posterior object if the statistic is sensitive to coherent redshift uncertainty, but it is more expensive to construct and store. For the BOSS CMASS-South tests in this paper, the difference between graphGP and KNN-field is smaller than the sample-variance error budget for the standard clustering statistics.

The report evaluates several statistic-level responses. The angular $w(\theta)$ response to shifting the missing-galaxy redshift prior is far below the realization scatter, as expected for a mostly angular statistic. Radial statistics are more sensitive. The completed $w_p(r_p)$ correction uncertainty is at the 0.2–0.4% level, and the real-data comparison with official weighted BOSS clustering is at the 1–5% level across the displayed projected and multipole statistics. For nonlinear summaries, the nearest-neighbor

CDF and counts-in-cells tests in Figure 7 show no large hidden failure from the local analogs or from the missing-redshift posterior. Users can repeat the same exercise for any proposed summary by running that summary on seeds $0, \dots, N_s - 1$.

6. SYSTEMATICS, MASKS, AND TRUSTWORTHINESS

The completion also produces diagnostic maps rather than only final catalogs. The redshift-failure coupling to the local target density is consistent with zero in the current BOSS CMASS-South product, $h = -0.06 \pm 0.06$, while the fiber-collision coupling is detected, $h = -0.24 \pm 0.03$. This is physically reasonable: collisions occur in denser angular environments. The completed density map correlates with the total-target density at 0.98, and the trustworthiness map has a median scatter of 0.13. These diagnostics identify where the completion is extrapolating most strongly.

The optional mask-inpainting product fills 542 interior holes totaling 16.6 deg^2 , about 0.9% of the footprint. The inpainting closure test is approximately 1% over 0.06° – 2° . We keep this step optional because different analyses may prefer to leave bright-star holes and small mask features explicit in the random catalog. The default product therefore corrects spectroscopic incompleteness and smooth imaging weights while leaving the official footprint visible.

7. DATA AND CODE AVAILABILITY

The analysis code is publicly available at <https://github.com/yiphey/graphGP-cosmology>, and the interactive completion report at <https://yiphey.github.io/graphGP-cosmology/completion.html>. The compact ECHOES posterior, the matching random catalog, and a dependency-light sampler are released with the report; a reproducible ensemble of completed catalogs is specified by the set of integer seeds used. The file contents, the minimal sampling interface, and the relevant source modules are documented in Appendix C.

8. CONCLUSIONS

We have presented ECHOES, an equal-weight completed-catalog posterior for BOSS DR12 CMASS-South. The product restores spectroscopically missing galaxies at real SDSS imaging positions, samples their redshifts from a data-driven posterior, and represents smooth imaging-systematic weights as local analog




FIG. 6.— Truth-known recovery of the projected correlation $w_p(r_p)$, plotted as the ratio to the input truth. *Left*: a real-BOSS-truth test in which the spectroscopic redshifts of a subset of CMASS-South galaxies are hidden and then completed. The incomplete observed catalog is biased low at small r_p (missing close pairs); the completed ensemble recovers the truth to 1–2% over $0.5 < r_p < 40 h^{-1} \text{Mpc}$. The oracle curve, which restores the hidden galaxies at their *true* redshifts, sets the achievable floor (ratio 0.997–1.005), so the residual is dominated by the redshift-assignment posterior rather than by the masking bookkeeping. *Right*: the same battery on independent MultiDark-Patchy mocks; recovery is faithful above $\simeq 1 h^{-1} \text{Mpc}$, and the sub-Mpc deficit is a known one-halo limitation of those mocks, confirmed by the real-BOSS-truth panel which recovers the small scales. The completion reproduces an unobserved truth, not merely the weighted bookkeeping of the observed catalog.




FIG. 7.— Higher-order, coincidence-sensitive validation of the completed catalogs against truth. *Left*: the nearest-neighbor (kNN) cumulative distribution functions for $k = 1, 2, 4$ — the distribution of distances from random points to the k -th nearest galaxy — which respond to duplicate points, missing close pairs, and redshift smearing. Completed and truth agree to a maximum $|\Delta \text{CDF}| \leq 0.0022$, and the absence of a spike at small separation confirms that the local-analog imaging-systematic step introduces no exact duplicates. *Right*: the counts-in-cells distribution in $8 h^{-1} \text{Mpc}$ spheres; the full one-point density distribution is recovered (mean 0.660 versus 0.663, variance-to-mean 2.37 versus 2.42, skewness 2.90 versus 2.94), not merely its second moment.




FIG. 8.— Calibration of the completion uncertainty against mock sample variance (MultiDark-Patchy). Per r_p bin, the realization-to-realization spread of the completed $w_p(r_p)$ (the completion uncertainty) is 0.17–0.44%, well below the 0.8–7.2% sample variance, so the completion adds only 6–36% of the sample-variance error in quadrature. Because the ensemble conditions on the single observed CMASS-South realization and resamples only the unobserved spectroscopic information, it is deliberately narrow and must be *added* to a cosmic-variance or simulation covariance, not used in place of it.

galaxies. It therefore changes the user-facing data interface: instead of asking every new summary statistic to understand the BOSS completeness weights, an analyst can run the statistic directly on completed point catalogs and propagate observational-completion uncertainty by resampling seeds.

The validation is deliberately conservative. The completed ensemble recovers truth-known $w_p(r_p)$ to the 1–2% level over the validated scales, matches official weighted BOSS clustering at the few-percent level, reproduces nearest-neighbor CDFs and counts-in-cells diagnostics, and has a completion uncertainty smaller than the CMASS-South sample variance. The graphGP route provides an independent, correlated posterior construction and agrees with the default KNN-field route within the current BOSS error budget.

The broader implication is practical. Survey collaborations have learned how to publish two-point-ready data products. The next generation of analyses will not be limited to two-point functions, and the next generation of surveys will have more complicated assignment, calibration, and selection functions. A ECHOES posterior over completed survey catalogs is a way to preserve the historical division of labor—survey teams ground observational corrections in the data, theory teams test statistics and models—while making the interface useful for a broad range of summaries. DESI, Euclid, Roman, CSST, and Rubin/LSST will all face versions of this problem. BOSS CMASS-South is small enough to solve carefully and mature enough to validate against the standard answer; it is therefore a useful prototype for how future public survey products can serve both standard and beyond-two-point cosmology.

APPENDIX

REPRODUCIBLE ALGORITHM

This appendix writes the default KNN-field completion as an implementation recipe.




FIG. 9.— The default KNN-field redshift engine versus the graphGP conditional-field route, on the real CMASS-South data. *Left:* projected and angular clustering of the two completions against the official w_c -weighted reference; graphGP, being a smooth field draw, sits closer to the weighted reference on larger projected scales (ratio near unity), while the KNN-field realization is modestly high in some bins. *Right:* truth-known recovery of $w_p(r_p)$ for both engines; the KNN-field engine is sharper on the sub-Mpc fiber-collision scale because it uses nearby observed redshifts directly, whereas graphGP is smoother and slightly closer on larger scales. Both recover the input clustering to the percent level; neither dominates, and the cheaper KNN-field route — which compresses to the public posterior — is the default.




FIG. 10.— The graphGP redshift assignment is invariant to the internal fiducial cosmology. *Left:* per-object missing redshifts assigned under a Planck and an Einstein-de Sitter fiducial agree with RMS ~ 0.001 and correlation 0.9996, far below the intrinsic assignment scatter and the photo- z uncertainty. *Right:* the recovered $w_p(r_p)$ differs by less than 0.1% between the two fiducials. The fiducial distance-redshift relation is a coordinate (gauge) choice absorbed by the data-measured kernel, not a cosmological prior.




FIG. 11.— Completion diagnostics. *Left:* the spectroscopic success fraction as a function of the local target overdensity for the two missing-target kinds. Redshift failures are consistent with no density coupling ($h = -0.06 \pm 0.06$), whereas fiber collisions are coupled ($h = -0.24 \pm 0.03$) because close pairs occur in denser environments; since the completion places collided galaxies at their true imaging positions, this coupling is reproduced by construction rather than modeled. *Right:* the trustworthiness map — the galaxy-count scatter across completion realizations per sky cell (median 0.13) — which marks where the redshift posterior leaves the most freedom and the completion is extrapolating most strongly.

1. Load the BOSS DR12 CMASS-South LSS catalog with photometry and the matching random catalog. Apply the cuts $RA < 28^\circ$ or $RA > 335^\circ$, $Dec > -6^\circ$, and $0.45 < z < 0.60$.
2. Convert SDSS MODELFLUX and EXTINCTION to magnitudes and colors. Train PhotoZKNN($k=100$) on finite observed galaxies using features $(g - r, r - i, i - z, i)$.
3. Load the CMASS target catalog. Build the missing list by greedily tying no-spectrum targets to WEIGHT_CP hosts within $62''$ and ZWARNING $\neq 0$ targets to WEIGHT_NOZ hosts within $15'$.
4. Measure the observed close-pair Δz pool from pairs of observed galaxies within $62''$ and symmetrize it.

5. For each missing target, find the $K = 150$ angular nearest observed spectroscopic galaxies. On a 256-point redshift grid, form the product of the local field KDE with $\sigma_f = 0.004$, the photo- z KDE with $\sigma_p = 0.02$, and, for collisions, the close-pair prior.
6. Store each normalized posterior as a 65-level inverse CDF. Quantize the observed redshifts and inverse-CDF values to `uint16`; store positions as `float32`, `WEIGHT_SYSTOT` as `float16`, and provenance as `int8`.
7. To draw seed s , invert each missing target’s CDF using uniform variates from `numpy.random.default_rng(s)`, add Gaussian redshift jitter $\sigma_f/2$, concatenate observed galaxies and restored targets, and add local imaging-systematic analogs with $1''$ angular jitter.
8. Measure statistics on many seeds. Use the seed mean as the completed survey statistic and the seed covariance as the completion covariance.

GRAPHGP IMPLEMENTATION DETAILS

The graphGP route uses the same observed BOSS catalog and missing target list but replaces the KNN line-of-sight density with a correlated posterior field. The implementation steps are:

1. Choose a fiducial distance-redshift relation for the internal field coordinate system and convert observed (α, δ, z) to $\mathbf{x}$.
2. Measure or tabulate the covariance $K_\xi(r)$ from the observed two-point function.
3. Build a sparse Vecchia neighbor graph over data and prediction points.
4. Generate a GP prior draw on that graph and condition it with Equation (12) using the Gaussianized Poisson likelihood.
5. Store $1 + \delta_{\text{post}}$ on a redshift-shell by HealPIX light cone.
6. For each missing target, evaluate Equation (13) on the same redshift grid used by the KNN-field sampler, then draw redshifts from the shared field realization.

The field posterior is not part of the default release because it is larger and more expensive, but it is implemented as a first-class alternative and was used for the comparisons in Section 4.6.

DATA PRODUCTS AND SAMPLER INTERFACE

The release in the report’s `docs/data/` directory contains the compact posterior `cmass_south_posterior.npz` (the fixed observed catalog plus an inverse-CDF redshift posterior for each restored target, tabulated at 65 quantile levels and quantized to `uint16` for a 2.04 MB total), the matching uniform-footprint randoms `cmass_south_randoms.npz`, and a NumPy-only sampler `draw_samples.py`. A completed catalog is drawn from an integer seed with

```
from draw_samples import load_package, draw
pkg = load_package("cmass_south_posterior.npz")
cat = draw(pkg, seed=0)
```

which returns the arrays `ra`, `dec`, `z`, `prov`, and `N`; a reproducible ensemble is the set of seeds used. The relevant source modules are `twopt_density/boss.py` (BOSS loader), `cmass_targets.py` (missing-target recovery), `photoz.py` (the color-space k NN photo- z), `observed_ls.py` (the KNN-field completion and the graphGP redshift interface), `posterior_sampler.py` (the compact inverse-CDF posterior), `density_field.py` (graphGP posterior-field sampling), and `clustering_corrfunc.py` (the projected and multipole clustering measurements).

This manuscript draft is based on the BOSS completion products in the `graphGP-cosmology` repository. The BOSS data products are from SDSS-III.

SOFTWARE

Corrfunc (M. Sinha & L. H. Garrison 2020), NumPy, SciPy, Astropy, Healpy, graphGP

REFERENCES

- | | |
|---|--|
| <p>Alam, S., Ata, M., Bailey, S., et al. 2017, The clustering of galaxies in the completed SDSS-III Baryon Oscillation Spectroscopic Survey: cosmological analysis of the DR12 galaxy sample, <i>MNRAS</i>, 470, 2617, doi: 10.1093/mnras/stx721</p> <p>Alcock, C., & Paczynski, B. 1979, An evolution free test for non-zero cosmological constant, <i>Nature</i>, 281, 358, doi: 10.1038/281358a0</p> | <p>Banerjee, A., & Abel, T. 2021, Nearest neighbour distributions: New statistical measures for cosmological clustering, <i>MNRAS</i>, 500, 5479, doi: 10.1093/mnras/staa3604</p> <p>Baugh, C. M., Gaztanaga, E., & Efstathiou, G. 1995, A comparison of the evolution of density fields in perturbation theory and numerical simulations - II. Counts-in-cells analysis, <i>MNRAS</i>, 274, 1049, doi: 10.1093/mnras/274.4.1049</p> |
|---|--|

- Biagetti, M., Calles, J., Castiblanco, L., Cole, A., & Noreña, J. 2022, Fisher forecasts for primordial non-Gaussianity from persistent homology, *J. Cosmology Astropart. Phys.*, 2022, 002, doi: 10.1088/1475-7516/2022/10/002
- Cheng, S., Ting, Y.-S., Ménard, B., & Bruna, J. 2020, A new approach to observational cosmology using the scattering transform, *MNRAS*, 499, 5902, doi: 10.1093/mnras/staa3165
- Cole, S., Percival, W. J., Peacock, J. A., et al. 2005, The 2dF Galaxy Redshift Survey: power-spectrum analysis of the final data set and cosmological implications, *MNRAS*, 362, 505, doi: 10.1111/j.1365-2966.2005.09318.x
- Colless, M., Dalton, G., Maddox, S., et al. 2001, The 2dF Galaxy Redshift Survey: spectra and redshifts, *MNRAS*, 328, 1039, doi: 10.1046/j.1365-8711.2001.04902.x
- CSST Collaboration, Gong, Y., Miao, H., et al. 2026, Introduction to the Chinese Space Station Survey Telescope (CSST), *Science China Physics, Mechanics, and Astronomy*, 69, 239501, doi: 10.1007/s11433-025-2809-0
- Davis, M., & Peebles, P. J. E. 1983, A survey of galaxy redshifts. V. The two-point position and velocity correlations., *ApJ*, 267, 465, doi: 10.1086/160884
- Dawson, K. S., Schlegel, D. J., Ahn, C. P., et al. 2013, The Baryon Oscillation Spectroscopic Survey of SDSS-III, *AJ*, 145, 10, doi: 10.1088/0004-6256/145/1/10
- de Lapparent, V., Geller, M. J., & Huchra, J. P. 1986, A Slice of the Universe, *ApJ*, 302, L1, doi: 10.1086/184625
- Einasto, J., Einasto, M., Gramann, M., & Saar, E. 1991, Structure and formation of superclusters. XIII - The void probability function, *MNRAS*, 248, 593, doi: 10.1093/mnras/248.4.593
- Eisenstein, D. J., Zehavi, I., Hogg, D. W., et al. 2005, Detection of the Baryon Acoustic Peak in the Large-Scale Correlation Function of SDSS Luminous Red Galaxies, *ApJ*, 633, 560, doi: 10.1086/466512
- Euclid Collaboration, Mellier, Y., Abdurro'uf, et al. 2025, Euclid: I. Overview of the Euclid mission, *A&A*, 697, A1, doi: 10.1051/0004-6361/202450810
- Feldman, H. A., Kaiser, N., & Peacock, J. A. 1994, Power-Spectrum Analysis of Three-dimensional Redshift Surveys, *ApJ*, 426, 23, doi: 10.1086/174036
- Friedrich, O., Uhlemann, C., Villaescusa-Navarro, F., et al. 2020, Primordial non-Gaussianity without tails - how to measure f_{NL} with the bulk of the density PDF, *MNRAS*, 498, 464, doi: 10.1093/mnras/staa2160
- Fry, J. N., & Gaztanaga, E. 1993, Biasing and Hierarchical Statistics in Large-Scale Structure, *ApJ*, 413, 447, doi: 10.1086/173015
- Guo, H., Zehavi, I., & Zheng, Z. 2012, A New Method to Correct for Fiber Collisions in Galaxy Two-point Statistics, *ApJ*, 756, 127, doi: 10.1088/0004-637X/756/2/127
- Hahn, C., Eickenberg, M., Ho, S., et al. 2023a, SIMBIG: mock challenge for a forward modeling approach to galaxy clustering, *J. Cosmology Astropart. Phys.*, 2023, 010, doi: 10.1088/1475-7516/2023/04/010
- Hahn, C., Eickenberg, M., Ho, S., et al. 2023b, A forward modeling approach to analyzing galaxy clustering with SIMBIG, *Proceedings of the National Academy of Science*, 120, e2218810120, doi: 10.1073/pnas.2218810120
- Hahn, C., Lemos, P., Parker, L., et al. 2024a, Cosmological constraints from non-Gaussian and nonlinear galaxy clustering using the SIMBIG inference framework, *Nature Astronomy*, 8, 1457, doi: 10.1038/s41550-024-02344-2
- Hahn, C., Eickenberg, M., Ho, S., et al. 2024b, Cosmological constraints from the nonlinear galaxy bispectrum, *Phys. Rev. D*, 109, 083534, doi: 10.1103/PhysRevD.109.083534
- Hamilton, A. J. S. 1993, Toward Better Ways to Measure the Galaxy Correlation Function, *ApJ*, 417, 19, doi: 10.1086/173288
- Harker, G., Cole, S., Helly, J., Frenk, C., & Jenkins, A. 2006, A marked correlation function analysis of halo formation times in the Millennium Simulation, *MNRAS*, 367, 1039, doi: 10.1111/j.1365-2966.2006.10022.x
- Hubble, E. 1934, The Distribution of Extra-Galactic Nebulae, *ApJ*, 79, 8, doi: 10.1086/143517
- Jasche, J., & Wandelt, B. D. 2013, Bayesian physical reconstruction of initial conditions from large-scale structure surveys, *MNRAS*, 432, 894, doi: 10.1093/mnras/stt449
- Kaiser, N. 1987, Clustering in real space and in redshift space, *MNRAS*, 227, 1, doi: 10.1093/mnras/227.1.1
- Katzfuss, M., Guinness, J., Gong, W., & Zilber, D. 2020, Vecchia Approximations of Gaussian-Process Predictions, *Journal of Agricultural, Biological and Environmental Statistics*, 25, 383, doi: 10.1007/s13253-020-00401-7
- Kitaura, F.-S., Rodríguez-Torres, S., Chuang, C.-H., et al. 2016, The clustering of galaxies in the SDSS-III Baryon Oscillation Spectroscopic Survey: mock galaxy catalogues for the BOSS Final Data Release, *MNRAS*, 456, 4156, doi: 10.1093/mnras/stv2826
- Lan, T.-W., Tojeiro, R., Armengaud, E., et al. 2023, The DESI Survey Validation: Results from Visual Inspection of Bright Galaxies, Luminous Red Galaxies, and Emission-line Galaxies, *ApJ*, 943, 68, doi: 10.3847/1538-4357/aca5fa
- Landy, S. D., & Szalay, A. S. 1993, Bias and Variance of Angular Correlation Functions, *ApJ*, 412, 64, doi: 10.1086/172900
- Lasker, J., Carnero Rosell, A., Myers, A. D., et al. 2025, Production of alternate realizations of DESI fiber assignment for unbiased clustering measurement in data and simulations, *J. Cosmology Astropart. Phys.*, 2025, 127, doi: 10.1088/1475-7516/2025/01/127
- Lemos, P., Parker, L., Hahn, C., et al. 2024, Field-level simulation-based inference of galaxy clustering with convolutional neural networks, *Phys. Rev. D*, 109, 083536, doi: 10.1103/PhysRevD.109.083536
- Massara, E., Villaescusa-Navarro, F., Ho, S., Dalal, N., & Spergel, D. N. 2021, Using the Marked Power Spectrum to Detect the Signature of Neutrinos in Large-Scale Structure, *Phys. Rev. Lett.*, 126, 011301, doi: 10.1103/PhysRevLett.126.011301
- Massara, E., Villaescusa-Navarro, F., Hahn, C., et al. 2023, Cosmological Information in the Marked Power Spectrum of the Galaxy Field, *ApJ*, 951, 70, doi: 10.3847/1538-4357/acd44d
- Massara, E., Hahn, C., Eickenberg, M., et al. 2025, Cosmological constraints using simulation-based inference of galaxy clustering with marked power spectra, *Phys. Rev. D*, 112, 083507, doi: 10.1103/914n-3pcx
- Miao, H., Gong, Y., Chen, X., et al. 2023, Cosmological constraint precision of photometric and spectroscopic multi-probe surveys of China Space Station Telescope (CSST), *MNRAS*, 519, 1132, doi: 10.1093/mnras/stac3583
- Mohammad, F. G., Percival, W. J., Seo, H.-J., et al. 2020, The completed SDSS-IV extended baryon oscillation spectroscopic survey: pairwise-inverse probability and angular correction for fibre collisions in clustering measurements, *MNRAS*, 498, 128, doi: 10.1093/mnras/staa2344
- Nguyen, N.-M., Schmidt, F., Tucci, B., Reinecke, M., & Kostić, A. 2024, How Much Information Can Be Extracted from Galaxy Clustering at the Field Level?, *Phys. Rev. Lett.*, 133, 221006, doi: 10.1103/PhysRevLett.133.221006
- Pinol, L., Cahn, R. N., Hand, N., Seljak, U., & White, M. 2017, Imprint of DESI fiber assignment on the anisotropic power spectrum of emission line galaxies, *J. Cosmology Astropart. Phys.*, 2017, 008, doi: 10.1088/1475-7516/2017/04/008
- Régalo-Saint Blancard, B., Hahn, C., Ho, S., et al. 2024, Galaxy clustering analysis with SimBIG and the wavelet scattering transform, *Phys. Rev. D*, 109, 083535, doi: 10.1103/PhysRevD.109.083535
- Reid, B., Ho, S., Padmanabhan, N., et al. 2016, SDSS-III Baryon Oscillation Spectroscopic Survey Data Release 12: galaxy target selection and large-scale structure catalogues, *MNRAS*, 455, 1553, doi: 10.1093/mnras/stv2382
- Ross, A. J., Beutler, F., Chuang, C.-H., et al. 2017, The clustering of galaxies in the completed SDSS-III Baryon Oscillation Spectroscopic Survey: observational systematics and baryon acoustic oscillations in the correlation function, *MNRAS*, 464, 1168, doi: 10.1093/mnras/stw2372
- Rybicki, G. B., & Press, W. H. 1992, Interpolation, Realization, and Reconstruction of Noisy, Irregularly Sampled Data, *ApJ*, 398, 169, doi: 10.1086/171845
- Satpathy, S., A C Croft, R., Ho, S., & Li, B. 2019, Measurement of marked correlation functions in SDSS-III Baryon Oscillation Spectroscopic Survey using LOWZ galaxies in Data Release 12, *MNRAS*, 484, 2148, doi: 10.1093/mnras/stz009
- Schäfer, F., Katzfuss, M., & Owahdi, H. 2021, Sparse Cholesky Factorization by Kullback-Leibler Minimization, *SIAM Journal on Scientific Computing*, 43, A2019, doi: 10.1137/20M1336254
- Schlegel, D. J., Finkbeiner, D. P., & Davis, M. 1998, Maps of Dust Infrared Emission for Use in Estimation of Reddening and Cosmic Microwave Background Radiation Foregrounds, *ApJ*, 500, 525, doi: 10.1086/305772
- Silber, J. H., Fagrellius, P., Fanning, K., et al. 2023, The Robotic Multiobject Focal Plane System of the Dark Energy Spectroscopic Instrument (DESI), *AJ*, 165, 9, doi: 10.3847/1538-3881/ac9ab1
- Sinha, M., & Garrison, L. H. 2020, CORRFUNC - a suite of blazing fast correlation functions on the CPU, *MNRAS*, 491, 3022, doi: 10.1093/mnras/stz3157
- Smee, S. A., Gunn, J. E., Uomoto, A., et al. 2013, The Multi-object, Fiber-fed Spectrographs for the Sloan Digital Sky Survey and the Baryon Oscillation Spectroscopic Survey, *AJ*, 146, 32, doi: 10.1088/0004-6256/146/2/32
- Szapudi, I., Meiksin, A., & Nichol, R. C. 1996, Higher Order Statistics from the Edinburgh/Durham Southern Galaxy Catalogue Survey. I. Counts in Cells, *ApJ*, 473, 15, doi: 10.1086/178123

- Szapudi, I., & Szalay, A. S. 1993, Higher Order Statistics of the Galaxy Distribution Using Generating Functions, *ApJ*, 408, 43, doi: 10.1086/172568
- ThBeyond-2pt Collaboration, Krause, E., Kobayashi, Y., et al. 2025, A Parameter-masked Mock Data Challenge for Beyond-two-point Galaxy Clustering Statistics, *ApJ*, 990, 99, doi: 10.3847/1538-4357/ad781d
- Uhlemann, C., Friedrich, O., Villaescusa-Navarro, F., Banerjee, A., & Codis, S. 2020, Fisher for complements: extracting cosmology and neutrino mass from the counts-in-cells PDF, *MNRAS*, 495, 4006, doi: 10.1093/mnras/staa1155
- Valogiannis, G., & Dvorkin, C. 2022, Towards an optimal estimation of cosmological parameters with the wavelet scattering transform, *Phys. Rev. D*, 105, 103534, doi: 10.1103/PhysRevD.105.103534
- Villaescusa-Navarro, F., Hahn, C., Massara, E., et al. 2020, The Quijote Simulations, *ApJS*, 250, 2, doi: 10.3847/1538-4365/ab9d82
- Wang, Y., Zhai, Z., Alavi, A., et al. 2022, The High Latitude Spectroscopic Survey on the Nancy Grace Roman Space Telescope, *ApJ*, 928, 1, doi: 10.3847/1538-4357/ac4973
- Wilding, G., Nevenzeel, K., van de Weygaert, R., et al. 2021, Persistent homology of the cosmic web - I. Hierarchical topology in Λ CDM cosmologies, *MNRAS*, 507, 2968, doi: 10.1093/mnras/stab2326
- Yip, J. H. T., Biagetti, M., Cole, A., Viswanathan, K., & Shiu, G. 2024, Cosmology with persistent homology: a Fisher forecast, *J. Cosmology Astropart. Phys.*, 2024, 034, doi: 10.1088/1475-7516/2024/09/034
- York, D. G., Adelman, J., Anderson, Jr., J. E., et al. 2000, The Sloan Digital Sky Survey: Technical Summary, *AJ*, 120, 1579, doi: 10.1086/301513
- Yuan, S., Abel, T., & Wechsler, R. H. 2024, Robust cosmological inference from non-linear scales with k-th nearest neighbour statistics, *MNRAS*, 527, 1993, doi: 10.1093/mnras/stad3359
- Zhan, H., & Tyson, J. A. 2018, Cosmology with the Large Synoptic Survey Telescope: an overview, *Reports on Progress in Physics*, 81, 066901, doi: 10.1088/1361-6633/aab1bd